

visualize

May 8, 2022

```
[1]: from sklearn.cluster import KMeans
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn import metrics
from sklearn.cluster import AgglomerativeClustering

from collections import Counter
import csv
```

```
[2]: # rawcount
from collections import Counter
import csv

def rawcount(filename, vocab, output):
    word_i = {}
    with open(vocab, 'r') as v:
        i = -1
        for word in v:
            if i == -1:
                i += 1
                continue
            word = word.strip()
            word_i[word.split(',')[1]] = i
            i += 1

    with open(output, 'w') as out:
        wr = csv.writer(out)

        with open(filename, 'r') as f:
            i = -1
            for line in f:
                # ignore the header
                if i == -1:
                    i += 1
                    continue
```

```

# remove \n
line = line.strip()

words = line.split(',')
p, seq = words[0], words[1:]
counter = Counter(seq)
p_rep = [0]*len(word_i)

for key, freq in dict(counter).items():
    p_rep[word_i[key]] = freq

wr.writerow([p]+p_rep)

```

1 Cluster Cohorts of Patients in Synthetic Data

```

[3]: convae_data_1 = pd.read_csv('synthetic/TRconvae-avg_vect_1.csv', header=None)
convae_data_1 = convae_data_1.drop(columns=[0])
convae_data_1.head()

```

```

[3]:
      1      2      3      4      5      6      7  \
0 -0.245492 -1.723636 -0.217432  1.412599  2.090927  3.460979  3.219759
1 -0.177317  0.581434 -0.233557 -0.819438 -0.856149  1.534644  0.902444
2 -0.086251  0.289699 -0.127992 -0.498970 -0.148464  1.362606  0.810772
3 -0.650934 -0.204592 -0.168182 -1.271896 -0.134454  1.193196  0.795659
4  0.338212  1.027026  0.182270 -1.075088 -0.242471  0.488513  1.398482

      8      9     10  ...     91     92     93     94  \
0  2.154350  1.517181  2.821212  ... -1.767284 -1.005675 -1.898477 -2.506755
1  0.126829  0.291722  0.764937  ... -0.202243  0.865371 -0.551938 -0.140060
2 -0.163232 -0.435123  1.222380  ... -0.430280  0.629954 -0.565389  0.090852
3  0.319215 -0.724538  1.937363  ... -0.066636 -0.058400 -1.086811 -0.447154
4 -0.588494  0.183105  0.659151  ... -0.729216  1.467192  0.488018 -0.057133

      95     96     97     98     99    100
0  2.391586  1.796674  0.014357 -1.714161 -1.905629 -2.658079
1  0.598072  0.185291  1.061913 -0.093946  0.092072  0.465387
2  0.130159  0.272310  0.473577 -0.101833 -0.001948  0.319315
3  0.839730  0.917877  0.798910 -0.865329 -0.244640 -0.172857
4 -0.086303 -0.434823  1.025620  0.104759  0.261447  0.593806

[5 rows x 100 columns]

```

```

[4]: convae_data_2 = pd.read_csv('synthetic/TRconvae-avg_vect_2.csv', header=None)
convae_data_2 = convae_data_2.drop(columns=[0])
convae_data_2.head()

```

```
[4]:
```

	1	2	3	4	5	6	7	\
0	3.460280	-1.165073	0.043982	-0.209492	0.691597	-0.023895	0.588930	
1	0.094355	-0.269476	-0.896734	-0.336747	-0.741758	0.211373	-0.197271	
2	-0.016761	-0.525858	-0.926781	0.099622	-1.112849	0.558424	0.014202	
3	0.121185	0.093348	-0.813224	-0.549337	-0.322482	0.427428	-0.452533	
4	3.246645	-0.045888	0.005351	0.497746	-0.444151	0.366363	1.445255	

	8	9	10	...	91	92	93	94	\
0	-0.065933	-0.062688	-2.780302	...	3.837300	1.916883	-0.023542	-0.230170	
1	-1.365462	-0.121160	0.344159	...	0.956573	0.295175	0.096312	-0.467132	
2	-1.164957	0.400407	0.595222	...	0.879721	0.042045	0.250709	-0.459051	
3	-0.515359	-0.242799	0.266890	...	0.112507	-0.083479	0.149594	-0.693066	
4	0.271684	-0.220651	-2.101012	...	3.391958	0.269805	0.052752	0.763482	

	95	96	97	98	99	100
0	-2.105041	3.084079	-1.048391	-1.242677	-1.625257	1.863790
1	-0.473703	0.175847	0.034651	-0.062616	0.048358	0.423669
2	-0.395686	0.296956	0.533112	-0.158044	-0.388362	0.519822
3	0.206137	-0.243773	0.206056	-0.331201	-0.546232	0.411023
4	-2.929876	1.655250	-0.830859	-1.148471	-2.996922	2.726214

[5 rows x 100 columns]

```
[5]: ## run only when first time generates the raw count of the synthetic data
# rawcount('synthetic/cohort-ehrseq.csv', 'synthetic/cohort-vocab.csv',
# ↪ 'synthetic/rawcount.csv')
```

```
[6]: rc_data = pd.read_csv('synthetic/rawcount.csv', header=None)
rc_data = rc_data.drop(columns=[0])
rc_data.head()
```

```
[6]:
```

	1	2	3	4	5	6	7	8	9	10	...	191	192	193	194	\
0	0	1	0	0	0	1	0	0	0	0	...	0	0	1	1	
1	0	0	0	0	0	1	0	0	1	0	...	0	0	0	0	
2	0	0	0	1	0	0	0	1	0	0	...	0	0	1	1	
3	0	0	0	0	0	0	0	0	1	0	...	0	0	0	1	
4	0	0	0	0	0	0	1	1	1	0	...	1	0	0	0	

	195	196	197	198	199	200
0	0	1	0	1	0	0
1	0	0	0	0	1	0
2	0	0	0	0	0	0
3	0	0	0	1	0	0
4	0	0	0	0	0	0

[5 rows x 200 columns]

```
[7]: rc_data = rc_data.values
      convae_data_1 = convae_data_1.values
      convae_data_2 = convae_data_2.values
```


```
[8]: def calculate_WSS(points, kmax):
      sse = []
      for k in range(1, kmax+1):
          kmeans = KMeans(n_clusters = k, random_state=1234).fit(points)
          centroids = kmeans.cluster_centers_
          pred_clusters = kmeans.predict(points)
          curr_sse = 0

          # calculate square of Euclidean distance of each point from its cluster
          ↪center and add to current WSS
          for i in range(len(points)):
              curr_center = centroids[pred_clusters[i]]
              curr_sse += (points[i, 0] - curr_center[0])**2 + (points[i, 1] -
          ↪curr_center[1])**2

          sse.append(curr_sse)
      return sse
```


```
[9]: wss_score = calculate_WSS(rc_data, 8)
      plt.plot(range(1, 9), wss_score)
      plt.title("WSS for Different k on RawCount Representation")
      plt.ylabel('WSS')
      plt.xlabel('k values')
```

```
[9]: Text(0.5, 0, 'k values')
```



```
[10]: wss_score = calculate_WSS(convae_data_1, 8)
plt.plot(range(1, 9), wss_score)
plt.title("WSS for Different k on ConvAE Representation")
plt.ylabel('WSS')
plt.xlabel('k values')
```

```
[10]: Text(0.5, 0, 'k values')
```



```
[11]: pca = PCA(2)

rc_data_2 = pca.fit_transform(rc_data)
convae_data_12 = pca.fit_transform(convae_data_1)
convae_data_22 = pca.fit_transform(convae_data_2)
```

1.1 k=2 clustering

```
[12]: cn = 2
km_rc = KMeans(cn).fit(rc_data)
km_convae1 = KMeans(cn).fit(convae_data_1)
km_convae2 = KMeans(cn).fit(convae_data_2)
```

```
[13]: print(metrics.silhouette_score(convae_data_1, km_convae1.labels_))
print(metrics.silhouette_score(convae_data_2, km_convae2.labels_))
print(metrics.silhouette_score(rc_data, km_rc.labels_))
```


```
0.6171488014869471
0.6804293786556501
0.08214529130164738
```

```
[14]: rc_label1 = rc_data_2[km_rc.labels_ == 0]

plt.scatter(rc_label1[:, 0], rc_label1[:, 1])
```

```
rc_label2 = rc_data_2[km_rc.labels_ == 1]
plt.scatter(rc_label2[:, 0], rc_label2[:, 1], color='red')
plt.title('K=2 Means Clustering for RawCount Representation')
plt.xlabel('PCA 1 feature')
plt.ylabel('PCA 2 feature')
```


[14]: Text(0, 0.5, 'PCA 2 feature')



```
[15]: convae_label1 = convae_data_12[km_convae1.labels_ == 0]

plt.scatter(convae_label1[:, 0], convae_label1[:, 1])
convae_label2 = convae_data_12[km_convae1.labels_ == 1]
plt.scatter(convae_label2[:, 0], convae_label2[:, 1], color='red')
plt.title('K=2 Means Clustering for ConvAE Representation')
plt.xlabel('PCA 1 feature')
plt.ylabel('PCA 2 feature')
```

[15]: Text(0, 0.5, 'PCA 2 feature')




1.2 k=5 clustering

```
[16]: cn = 5
      km_rc = KMeans(cn).fit(rc_data)
      km_convae_12 = KMeans(cn).fit(convae_data_1)
      km_convae_22 = KMeans(cn).fit(convae_data_2)


[17]: colors = ['red', 'blue', 'orange', 'black', 'green']
      for label in range(5):
          rc_label = rc_data_2[km_rc.labels_ == label]
          plt.scatter(rc_label[:, 0], rc_label[:, 1], color=colors[label])

      plt.title('K=5 means clustering for RawCount representation')
      plt.xlabel('PCA 1 feature')
      plt.ylabel('PCA 2 feature')
      plt.show()
```

```
[18]: colors = ['red', 'blue', 'orange', 'black', 'green']
      for label in range(5):
          convae_label = convae_data_12[km_convae_12.labels_ == label]
          plt.scatter(convae_label[:, 0], convae_label[:, 1], color=colors[label])

      plt.title('K=5 means clustering for ConvAE representation')
      plt.xlabel('PCA 1 feature')
      plt.ylabel('PCA 2 feature')
      plt.show()
```



2 Suicide Risk Assessment based on Reddit Posts

2.1 Data Preprocess(no need to run this when evaluating the model)

```
[19]: import re
import collections

data = suicide_data = pd.read_csv('500_Reddit_users_posts_labels.csv')

posts = []
vocabulary = {}

for userpost in data.Post.values:
    posts.append([])
    for post in userpost[2:-2].split('\n'):
        posts[-1].append(re.sub('[^a-zA-Z0-9 \n]', '', post).lower())

        for word in posts[-1][-1].split(' '):
            if word not in vocabulary:
                vocabulary[word] = len(vocabulary)
posts[-1] = ''.join(posts[-1])
```

```

train_idx = np.random.choice(len(posts), int(0.7*len(posts)), replace=False)
test_idx = np.array(list(set(range(len(posts))).difference(set(train_idx))))
train_posts = np.array(posts)[train_idx]
test_posts = np.array(posts)[test_idx]

with open('cohort-ehrseq.csv', 'w') as train_out:
    train_out.write('MRN,EHRseq\n')

    for i, p in enumerate(train_posts):
        word_list = p.split(' ')
        word_emb = [str(vocabulary[w]) for w in word_list]
        word_str = ','.join(word_emb)
        train_out.write(f'pat_{train_idx[i]},{word_str}\n')

with open('cohort_test-ehrseq.csv', 'w') as test_out:
    test_out.write('MRN,EHRseq\n')

    for i, p in enumerate(test_posts):
        word_list = p.split(' ')
        word_emb = [str(vocabulary[w]) for w in word_list]
        word_str = ','.join(word_emb)
        test_out.write(f'pat_{test_idx[i]},{word_str}\n')

with open('cohort-vocab.csv', 'w') as vocab_out:
    vocab_out.write('LABEL,CODE\n')

    for key in vocabulary:
        vocab_out.write(f'{key},{vocabulary[key]}\n')

np.save(
    open('train_idx.txt', 'wb'),
    train_idx
)
np.save(
    open('test_idx.txt', 'wb'),
    test_idx
)

```

2.2 Visualize Clustering

```

[20]: convae_data_1 = pd.read_csv('suicide/TRconvae-avg_vect_1.csv', header=None)
convae_data_1 = convae_data_1.drop(columns=[0])
convae_data_1.head()

```

```

[20]:
      1      2      3      4      5      6      7  \
0  4.792002 -4.182917  0.504329  1.034935  6.722288  2.965201 -7.597964
1  3.031826 -2.228876  0.723921  1.243128  0.335348  0.966756 -2.259586

```

2	4.347315	0.317548	9.783994	1.039316	-6.086713	-3.634538	-2.878214
3	4.722687	1.320996	12.259055	0.045357	-7.071722	-4.573728	-3.511061
4	2.957799	-2.804701	-1.197343	1.118323	2.548878	1.793160	-3.890705

	8	9	10	...	91	92	93	94	\
0	6.288152	1.327902	-3.691782	...	-4.469840	-0.809733	2.743039	2.791558	
1	2.042837	1.677121	-2.544549	...	-2.905806	-1.389448	1.969885	0.689893	
2	2.623119	0.926128	-3.325168	...	-2.733623	-2.882202	4.613301	-1.638416	
3	2.502617	0.535087	-2.168540	...	-2.568284	-2.593325	6.170129	-3.251006	
4	1.929259	1.277872	-2.188694	...	-2.946431	-0.506851	1.684679	1.369734	

	95	96	97	98	99	100
0	3.434430	5.926430	-1.437946	-2.517109	5.947384	3.499781
1	-0.540738	2.790236	0.998298	-2.344237	3.917216	1.695267
2	-0.787181	1.801085	2.008439	-2.242496	3.588685	0.511575
3	-0.783794	1.943220	1.631231	-3.375519	3.806715	-0.969864
4	0.733448	3.910773	0.118717	-2.856843	4.061832	1.492922

[5 rows x 100 columns]

```
[21]: convae_data_2 = pd.read_csv('suicide/TRconvae-avg_vect_2.csv', header=None)
convae_data_2 = convae_data_2.drop(columns=[0])
convae_data_2.head()
```

```
[21]:
```

	1	2	3	4	5	6	7	\
0	-0.265760	1.819030	2.114592	-1.503902	1.709242	1.080806	2.508768	
1	-0.111605	2.074826	2.444105	-1.740146	1.859789	0.832969	2.819070	
2	-1.081815	4.426412	3.382702	-3.064773	1.503283	2.644223	3.081339	
3	-1.153381	2.246483	1.979828	-1.829293	1.712191	1.823514	2.249827	
4	-0.747093	3.698134	2.605899	-3.106969	1.862821	2.313798	2.278987	

	8	9	10	...	91	92	93	94	\
0	2.006261	-2.160024	1.149742	...	-0.761870	0.483298	2.443803	-0.542755	
1	2.144344	-2.465429	1.165465	...	-0.809549	0.820165	2.769438	-0.086744	
2	1.536571	-4.454443	0.523732	...	-0.469871	0.311442	2.297619	0.251516	
3	1.622034	-2.578066	0.850274	...	-0.643407	-0.029906	1.923287	-0.989684	
4	1.505187	-3.374255	0.680390	...	-1.077267	-0.205817	1.848762	0.063171	

	95	96	97	98	99	100
0	0.534815	-1.088351	-2.180832	-2.487342	-1.508296	-1.993607
1	0.541063	-1.366083	-2.689430	-2.858235	-1.534366	-2.045888
2	-0.340333	-2.125811	-5.249793	-2.501698	-1.029516	-1.661276
3	0.266016	-1.610323	-2.569872	-2.158781	-1.231015	-1.765235
4	-0.293556	-1.952785	-4.631961	-2.074854	-1.011812	-1.770192

[5 rows x 100 columns]

```
[22]: rc_data = pd.read_csv('suicide/suicide_rawcount.csv', header=None)
rc_data = rc_data.drop(columns=[0])
rc_data.head()
```

```
[22]:
```

	1	2	3	4	5	6	7	8	9	10	...	\
0	6	7	22	0	2	38	3	15	0	10	...	
1	0	1	2	0	0	1	0	0	0	0	...	
2	4	4	27	0	0	26	2	10	0	18	...	
3	0	11	42	0	0	76	1	26	0	18	...	
4	0	1	2	0	0	8	0	1	0	2	...	

	22920	22921	22922	22923	22924	22925	22926	22927	22928	22929
0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0

[5 rows x 22929 columns]

```
[23]: convae_data1 = convae_data_1.values
convae_data2 = convae_data_2.values
rc_data = rc_data.values
```

```
[24]: pca = PCA(2)

rc_data_2 = pca.fit_transform(rc_data)
convae_data_12 = pca.fit_transform(convae_data1)
convae_data_22 = pca.fit_transform(convae_data2)
```

```
[25]: # generate labels
suicide_label_map = {'Supportive':0, 'Indicator':1, 'Ideation':2, 'Behavior':3,
↳ 'Attempt':4}
mapLabel = np.vectorize(lambda x: suicide_label_map[x])

suicide_data = pd.read_csv('suicide/500_Reddit_users_posts_labels.csv')
suicide_label = mapLabel(suicide_data.Label)

train_idx = np.load('suicide/train_idx.txt')
test_idx = np.load('suicide/test_idx.txt')
```

```
[26]: suicide_train_label = suicide_label[train_idx]
```


```
[27]: colors = ['green', 'blue', 'orange', 'red', 'black']
labels = ['Supportive', 'Indicator', 'Ideation', 'Behavior', 'Attempt']
for label in range(5):
    convae_label1 = convae_data_12[suicide_train_label == label]
```

```

plt.scatter(convae_labeli[:, 0], convae_labeli[:, 1], color=colors[label],
↪label=labels[label])
plt.title('Clustering for ConvAE(1-layer) Representation')
plt.xlabel('PCA 1 feature')
plt.ylabel('PCA 2 feature')
plt.legend()

```

[27]: <matplotlib.legend.Legend at 0x137caca30>




```

[28]: colors = ['green', 'blue', 'orange', 'red', 'black']
labels = ['Supportive', 'Indicator', 'Ideation', 'Behavior', 'Attempt']
for label in range(5):
    convae_labeli = convae_data_22[suicide_train_label == label]
    plt.scatter(convae_labeli[:, 0], convae_labeli[:, 1], color=colors[label],
↪label=labels[label])
plt.title('Clustering for ConvAE(2-layer) Representation')
plt.xlabel('PCA 1 feature')
plt.ylabel('PCA 2 feature')
plt.legend()

```


[28]: <matplotlib.legend.Legend at 0x13853d100>



```
[29]: import matplotlib.pyplot as plt

colors = ['green', 'blue', 'orange', 'red', 'black']
for label in range(5):
    rc_labeli = rc_data_2[suicide_train_label == label]
    plt.scatter(rc_labeli[:, 0], rc_labeli[:, 1], color=colors[label],
        ↪label=labels[label])
plt.title('Clustering for RawCount Representation')
plt.xlabel('PCA 1 feature')
plt.ylabel('PCA 2 feature')
plt.legend()
```

```
[29]: <matplotlib.legend.Legend at 0x1368afca0>
```



2.3 Hierarchical Clustering

```
[30]: def purity_score(y_true, y_pred):
        contingency_matrix = metrics.cluster.contingency_matrix(y_true, y_pred)
        return np.sum(np.amax(contingency_matrix, axis=0)) / np.
        ↪sum(contingency_matrix)
```

```
[31]: rc_hierarchical = AgglomerativeClustering(n_clusters=5).fit(rc_data)
        convae1_hierarchical = AgglomerativeClustering(n_clusters=5).fit(convae_data_1)
        convae2_hierarchical = AgglomerativeClustering(n_clusters=5).fit(convae_data_2)
```

```
[32]: print(f'Purity for ConvAE with 1-layer CNN is_
        ↪{purity_score(convae1_hierarchical.labels_, suicide_train_label)}')
        print(f'Purity for ConvAE with 2-layer CNN is_
        ↪{purity_score(convae2_hierarchical.labels_, suicide_train_label)}')
```

Purity for ConvAE with 1-layer CNN is 0.62

Purity for ConvAE with 2-layer CNN is 0.84

```
[33]: print(f'Entropy for ConvAE with 1-layer CNN is {metrics.
        ↪homogeneity_score(convae1_hierarchical.labels_, suicide_train_label)}')
        print(f'Entropy for ConvAE with 2-layer CNN is {metrics.
        ↪homogeneity_score(convae2_hierarchical.labels_, suicide_train_label)}')
```



```
Entropy for ConvAE with 1-layer CNN is 0.0219725368773614  
Entropy for ConvAE with 2-layer CNN is 0.03837413298317611
```

```
[ ]:
```